

RETRIEVAL LATENCY DECOMPOSITION

End-to-end retrieval latency is decomposed so optimization targets the actual tail contributor instead of guessing.

